



## **School of Electrical and Computer Engineering**

### **CptS223: Advanced Data Structures C/C++**

**Fall 2019**

#### **Programming Assignment 5 (PA5)**

#### **Priority Queues**

**Due: Friday 11/15/2019 @ 11:59pm**

### **0. Instructions**

All source code must be in C++. You are required to use the Unix environment. Your submission should accompany a report that explains your work. Your code will be tested on EECS servers. Thus, you are strongly encouraged to test (or even write from the beginning) your code on EECS servers before submitting them. If you do not have an EECS account, you should get one immediately. The EECS IT team will assist you with setting up your EECS account. Look for IT on this page for contact information of the IT personnel:

<https://school.eecs.wsu.edu/people/staff/>

If you have not logged into EECS Gitlab, you should do so immediately. Information about how to install Git on your computer and how to upload your package (code + report) to the EECS Git can be found under “**Programming Assignments**” section of the course web page. All submission files should be submitted through Git. Note that submissions through any other means for example by emailing your package to the instructor/TA/TAs will be discarded and will NOT be graded. So please follow the above instructions carefully.

You can find the course web page following this link:

<http://eps1.eecs.wsu.edu/teaching/>

For this programming assignment, you have the option of either working as a TEAM if you prefer, or as an individual. A TEAM should contain at most 2 students including yourself. Grading will NOT differentiate between the two - i.e., if you are working as a team of two people, then the same grade will be applied to both of you.

If working as a team, both of you should submit the same package. However, you will need to clearly mention team members names in your report. If working as a team, also remember to include both your last names in the file name of your report.

In the rest of this document, the term “YOU” is used either to refer to you (if you are working individually) or your team (if you work as a team).

The final material for submission should be entirely written by you. If you decide to consult with others or refer materials online, you MUST give due credits to these sources (people, books, webpages, etc.) by listing them in the report that accompanies your submission. Note that no points will be deducted for referencing these sources. However, your discussion/consultation should be limited to the initial design level (if any). Sharing or even showing your source code/assignment solution verbiage to anyone else in the class, or direct reproduction of source code/verbiage from online resources, will all be considered plagiarism, and will therefore be awarded **ZERO** points and subject to the WSU Academic Dishonesty policy. (Reproducing from the Weiss textbook is an exception to this rule, and such reproduction is encouraged wherever possible.)

Late submission policy: A late penalty of 10% will be assessed for late submissions within the next 24 hours (i.e., until midnight of the next day after the submission deadline). After this one-time late submission, no submissions will be accepted.

## 1. Introduction

You will be primarily using priority queues for this programming assignment. Among priority queues, you are free to pick and choose the most appropriate/efficient heap implementation. Your design can also use any other data structures that we have discussed in class, if the use of such a data structure could be justified from a design perspective. However, you are required to choose a simple data structure wherever possible. An unnecessarily overcomplicated implementation (example: using a heap where a simple vector would suffice) will lead to reduction in points. The idea is to let you make the design choice and identify the most suitable data structure for each component in the system in order to effectively address performance and functional tradeoffs. In your report, please discuss your design and your design choices.

Use of any appropriate STL containers is highly encouraged whenever possible (priority queue, set, map, list, vector, etc.). Recall that STL's "priority queue" implements a binary heap in the form of a *MaxHeap*. You can easily change it into a *MinHeap* (as shown in the book, Figure 6.57). You are also free to re-use/reproduce any source code that is provided in the Weiss textbook. Most of the source codes are available from the following website:

[https://users.cs.fiu.edu/~weiss/dsaa\\_c++4/code/](https://users.cs.fiu.edu/~weiss/dsaa_c++4/code/)

## 2. Problem Statement

Your organization has purchased a new parallel computer (or "cluster") which has ' $p$ ' processors. Your task is to design and implement a simple shortest-job-first scheduler that allows multiple users to access the cluster at the same time, as per the specifications in Section 3 of this document. Figure 1 is a conceptual illustration of the different components in the system, how they interact, and the necessary data containers you will have to implement. Refer to this figure while reading the rest of this document.

### 3. Problem Specification

The scheduler performs all required functions at regular timesteps, called “ticks”. The tick in other words, is the pulse of the scheduler. During each tick, various events could happen, and the scheduler’s responsibility is to address all these events before moving on to the next tick. Because of this, each tick could take variable amount of time to complete, and for design simplicity you can assume that there is enough real time within each tick to complete all required tasks. The various events and the respective action that is required to be carried out by the scheduler are outlined below. (Refer Figures 1 & 2 as you read this.)

Init condition: When the algorithm starts, the system is initialized with  $p$  processors in the free pool. Figure 2 shows the sequence of steps that the scheduler should perform during each tick as a pseudocode.

**1. Event:** The user inputs at most one job at the start of any tick. Each job arrives to the scheduler in the following format:  $\langle job\_description, n\_procs, n\_ticks \rangle$ . If the *job\_description* field is set to “NULL” then that means that the user does not have any job to insert during this clock tick.

**Action:** For each job submission, the system assigns the job a new, unused integer id (referred to as “*job\_id*”) and then calls the *InsertJob()* function with the following parameters:

$\langle job\_id, job\_description, n\_procs, n\_ticks \rangle$ , where:

- a. “*job\_description*” contains details on what command to execute to run the job (i.e., “ $\langle program\ name \rangle \langle its\ arguments \rangle$ ”). It is a string data type. If the *job\_description* is set to “NULL” then it means that there are no jobs to be inserted in this clock tick;
- b. “*n\_procs*” is the number of processors that this job needs to run on;

- c. “*n\_ticks*” is the number of ticks that this job is going to need on each of the *n\_procs* processors. Note that this number is the equivalent of the job’s estimated running time, from start to finish, as the job will be launched on all *n\_procs* processors at the same time to run in parallel for *n\_ticks* number of ticks.

**Note:** The notation *InsertJob()* does not mean there are no arguments to this function. You need to decide what arguments each function signature should have. I have purposefully left that choice to you as part of design. Same applies to all the other functions.

The *InsertJob()* function first checks if  $(0 < n\_procs \leq p)$  and  $(n\_ticks > 0)$ . If so, it inserts the job into a “wait queue”. Otherwise, a job submission error is raised with an appropriate message.

**2. Event:** The job wait queue is not empty.

**Action:** Among all the jobs in the wait queue, find the job that is expected to take the least time to complete (*FindShortest()*) – i.e., a job with the minimum value for *n\_ticks*. Let  $J_i$  denote the selected job, and  $p_i$  denote its *n\_procs*. Now, check if there are at least  $p_i$  processors currently available in the processor “free pool” (*CheckAvailability()*). If not, then leave the wait queue intact and return. Otherwise:

- a. remove  $J_i$  from the wait queue (*DeleteShortest()*), and
- b. remove  $p_i$  number of arbitrary processors from the free pool and “assign” them to that job (*RunJob()*), so that they can start executing the job from the next tick. When a job goes to the running state, along with it a countdown “timer” is started; initialize this to *n\_ticks*.

If after the first job allocation, there are still more jobs in the wait queue that have sufficient available processors to run on, then iteratively keep removing all the next shortest jobs until no more processors are available for the next shortest job, or the wait queue becomes empty.

**Note:** remember your code does not have to concern itself about running the actual job. You can safely assume that will be internally taken care of by the individual processors that you assigned to run the job. Your responsibility is to make sure that the scheduler is aware of how many processors are now available in the free pool after the allocation and how many processors this new job that got moved to the run queue is using

**3. Housekeeping Event:** The running job queue is not empty.

**Action:** The job timer for ALL running jobs is decremented by one (*DecrementTimer()*).

**4. Housekeeping Event:** A job that was in the running state completed during this tick – i.e., its timer has reached a value of zero after decrement.

**Action:** The scheduler should release all the processors attached to this job, back into the free pool so that they become available for other jobs during later ticks (*ReLeaseProcs()*).

## 4. API and Testing

Follow the below instructions to test your code.

- Write a separate `TestDriver.cpp` program which first reads an input file, in which a series of job inserts are mentioned, one line per tick. An example input is shown below:

J1	8	10
J2	2	1
J3	12	12
J4	10	2
J5	5	8
J6	4	2
J7	4	6
J8	2	5
J9	4	3
J10	6	2

Your code should iterate through these lines, one tick per each line, and adding any more iterations (i.e., ticks) as necessary until the Wait Queue becomes empty. Within each iteration, the current tick number should be displayed (starting at serial number '1'). Then the code should call the *Tick()* function (as specified in the pseudocode in Figure 2. Define the *Tick()* function as a member function of your Scheduler class.

- At the end of the *Tick()* function, print a system-defined `job_id` number unique to any job submitted during this tick or an error message upon failure to insert. Also print all jobs which got allocated to the running queue in this tick, and also any job or jobs that completed during this tick. Each print message should include the job description, number of processors it needs, and the number of ticks it needed.

- The test program should terminate if the user inputs “exit” for the job descriptor in the standard input.

## 5. Design Document

Along with your source code, also submit a 1-2 pages design document (as Word Doc or PDF - no other formats allowed).

In the first part of this document, provide a figure that best describes your design for the scheduler system. It should basically be a modified version of Figure 1, in which you should state what data structures/STL containers you used to implement the three data containers (wait queue, running job queue, free pool).

In the second part, make a table to simply state the worst-case run-time complexities in your implementation for each of the event functions: ***InsertJob***, ***FindShortest***, ***DeleteShortest***, ***CheckAvailability***, ***RunJob***, ***DecrementTimer***, ***ReleaseProcs***. No need for justifications. Just write the answers using the Theta (or Big-Oh) notation. You can use ‘*n*’ to denote the maximum number of jobs (waiting or running) at any given point of time, and ‘*p*’ to denote the total number of processors (free or busy).

In the third part, write a short paragraph about what you think are the main shortcomings/bottlenecks of this shortest-job-first strategy, from both a performance and functionality point of view. Support your claims with brief but compelling rationale.



## 6. Grading Rubrics

This assignment is mainly about designing efficient algorithms and implementations for the Scheduler class. There is not really an extensive empirical component that involves timing and such but you need to pay attention to design in a manner that is expected to keep the runtimes really low for even large input cases. So for grading this assignment, we will primarily look at how well you have designed your algorithm behind the implementation, how well you have implemented the code, and whether your code tests correctly.

### Algorithm Design, Coding, and Testing (70 pts):

- (30 pts): Is the algorithm design efficient? That is, does the algorithm deploy the appropriate kind of data structures and use them effectively in a way that is expected to deliver the best possible performance both in terms of runtime and memory?
- (20 pts): Is the code implemented in an efficient way? i.e., are there parts in the code that appear redundant, or implemented in ways that can be easily improved? Does the code conform to good coding practices of Objected Oriented programming?
- (10 pts): Is the code documented well and generally easy to read (with helpful comments and pointers)?
- (10 pts): Does the code compile and run successfully on a couple of test cases? Also does the output for the sample test case provided look right?

### Report (30 pts):

- (20 pts): Is the algorithm described well in the report? Are all important details effectively presented, or does the TA have to dig up some critical information (such as which STL classes were used to implement the Scheduler class or how each of its main methods (both housekeeping and even-driven, see Figure 1) was implemented using different data structures) from the code?
- (10 pts): Is the basis of the choice of different data structures well justified/supported with the presentation of the expected time and memory complexities?

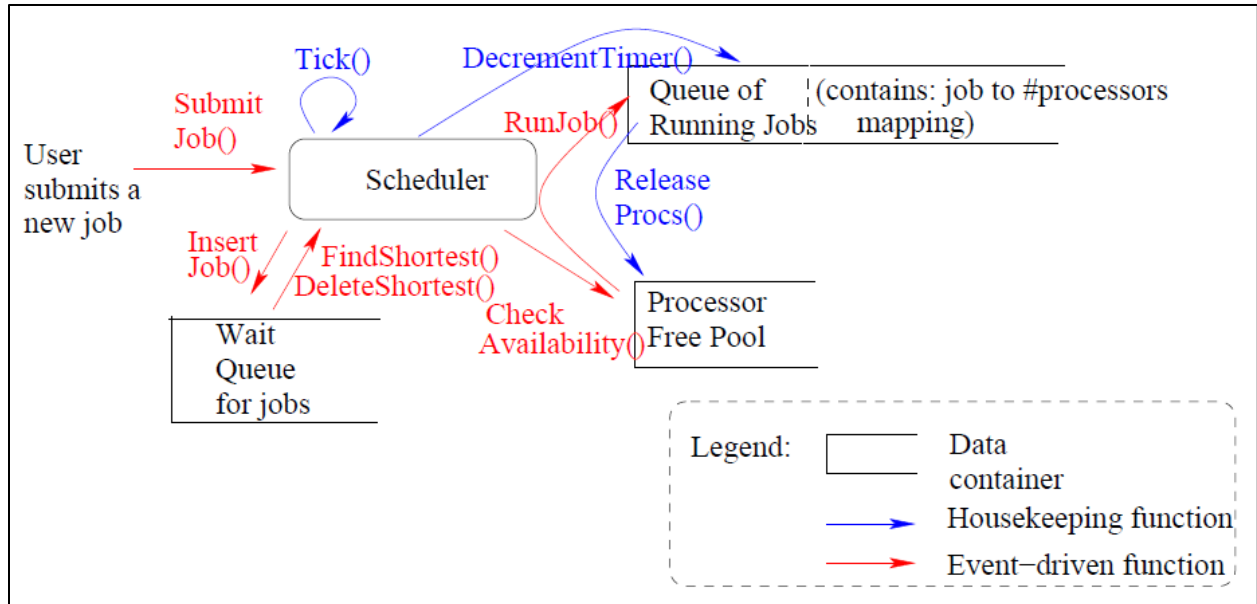


Figure 1. Organization of the scheduler system for a parallel cluster that has a total of  $p$  processors.

#### Algorithm Pseudocode for the scheduler' Tick procedure

Tick()

0. Prompt user for submitting a new job.
1. Insert the job into the wait queue (if the user submitted a valid job);
2. Decrement timer for all running jobs in the running queue;
3. Release all the processors corresponding to the completed/expired jobs from the running queue to the free pool. A job is said to have "expired" if its timer becomes zero during this tick. Also, any expired job should be automatically removed from the running job queue;
4. Find the next shortest job  $\langle J_i, p_i, t_i \rangle$  from the wait queue and check if sufficient number of processors are available in the free pool to run it. If so:
  - a. Remove job  $J_i$  from the wait queue;
  - b. Insert job  $J_i$  into the running queue and assign  $p_i$  processors (from free pool) to that job; If after this first job allocation, there are still more jobs in the wait queue that have sufficient available processors to run on, then iteratively keep removing all the next shortest jobs until no more processors are available for the next shortest job, or the wait queue becomes empty.

Figure 2. Sequence of steps to be carried out during each Tick() of the scheduler.